

Jan-28 Lecture

Multimedia Data

- From a networking perspective, what is different about multimedia data?

Multimedia Data

- From a networking perspective, what is different about multimedia data?
 - In some cases, nothing... if it is just a file transfer, then there are no particular additional requirements
 - Only when data is “real-time” is there any point to study
 - “soft” real-time: used/consumed as soon as it is received, the faster the better
 - “hard” real-time: if it doesn’t make it by a deadline, bad things happen

Multimedia Data

- From a networking perspective, what is different about multimedia data?
 - In some cases, nothing... if it is just a file transfer, then there are no particular additional requirements
 - Only when data is “real-time” is there any point to study
 - “soft” real-time: used/consumed as soon as it is received, the faster the better
 - “hard” real-time: if it doesn’t make it by a deadline, bad things happen
- For improved delivery, there are two choices:
 1. Implement support “in the network” (e.g., quality of service)
 - Better than best-effort
 2. Implement support “on top of the network”
 - Transport-layer support, session-layer support, application-layer support

Multimedia Data

- Support “in the network” (e.g., quality of service)
 - Mark packets according to their data type
 - Queue/forward different types of packets differently
 - Create a (logical) connection-based system and reserve more/less resources for different types of packets

Multimedia Data

- Support “in the network” (e.g., quality of service)
 - Mark packets according to their data type
 - Queue/forward different types of packets differently
 - Create a (logical) connection-based system and reserve more/less resources for different types of packets
- Implement support “on top of the network”
 - Transport layer: TCP or UDP
 - Session layer: UDP+, RTP, RTSP, UPnP, DLNA
 - Application layer: Proprietary solutions and/or at least application-specific solutions

When Streaming Was a Thing

- Links were just able to support bandwidth requirements
 - End-to-end path had to support delivery rate at close to the consumption rate
 - If the path couldn't keep up, playout would stop
 - Could build in buffering at the beginning to avoid interruptions
 - Works okay for non-real-time, but not for VoIP

When Streaming Was a Thing

- Links were just able to support bandwidth requirements
 - End-to-end path had to support delivery rate at close to the consumption rate
 - If the path couldn't keep up, playout would stop
 - Could build in buffering at the beginning to avoid interruptions
 - Works okay for non-real-time, but not for VoIP
- In either situation, it was critical *not* to stop the transmission of data
 - So no TCP
 - Was fine since video could be made robust against errors
- The big challenge was regulating the flow of data from source to decoder

When Streaming Was no Longer a Thing

- Links became faster—faster than the rate at which high quality video required more bandwidth
 - The kinds of delays that caused playout interruptions diminished
 - Data could be delivered faster than it was played out
 - TCP became viable again
 - Especially since UDP was used for DoS attacks
 - Especially since UDP was not nice to TCP
 - Sort of still important to regulate the flow of data, but not really
 - Buffering/Memory was increasingly cheap
 - But still some motivation not just to deliver all of the data
 - Still applications like VoIP—data is not available until it is created
 - But here just use minimal buffers, primarily for jitter and TCP behaviors
 - TCP behaviors become increasingly rare in high bandwidth, low congestion environments

Summary of Influences on Throughput

- Rate at which data is made available
 - A factor for live audio/video...
 - ...but not for data drawn from a file
- TCP behaviors
 - Certainly slow start, but that doesn't last long
 - Can be a factor if link quality is an issue
- Application layer control
 - Important to have some control of data flow
 - Especially if link is fast and data is coming from a file

Streaming v. Downloading

- Download: fetch/retrieve the complete file
- Streaming: fetch/retrieve data “just-in-time”

Responsiveness

Less
Storage

More
Storage

More
Responsive

Less
Responsive

Storage Requirements

- Lots of techniques to create hybrids of these options

Hybrid Techniques

- Download: fetch/retrieve the complete file
- Streaming: fetch/retrieve data “just-in-time”
- Progressive download (or buffered streaming)
 - Download first part of the content first
 - Same as streaming with a significant buffer
 - [Q]: How much to download/buffer before starting?
 - Wait a fixed amount of time
 - Dynamically determine based on download and playout rates
 - Also have to deal with buffer management

Back to Data Type

- The type of application has an impact on the techniques that can be used
- Most demanding application is **video conferencing**
 - Needs to be interactive and has high bandwidth requirements
 - **Voice-over-IP is a close second** (same interactive requirements but the bandwidth requirements aren't as high)
 - Also worth mentioning **MMOG**
- Most other applications in the Internet only have the timing requirement of “I want it as fast as possible”
- Worth talking about errors and resiliency

Sources of Delay in Interactive Voice

- From the time a word is spoken to the time the sound wave is received at the destination user's ear, what are the major sources of delay?

Sources of Delay in Interactive Voice

- From the time a word is spoken to the time the sound wave is received at the destination user's ear, what are the major sources of delay?
 1. Encoding/compression delay
 2. Packetization delay
 3. Propagation delay
 4. Decoding/decompression delay

Sources of Delay in Interactive Voice

- From the time a word is spoken to the time the sound wave is received at the destination user's ear, what are the major sources of delay?
 1. Encoding/compression delay
 2. Packetization delay
 3. Propagation delay
 4. Decoding/decompression delay
- Another serious concern is jitter
 - In interactive applications, not much you can do about jitter
 - In less stringent streaming applications, can use buffering

Now on to Audio

- Focus today on audio encoding and compression
 - Wednesday on video encoding and compression

Audio Signals

- All sound is an analog signal
 - What we say and what we hear are continuous wave forms
- *Goal:* reduce the analog signal into a series of 1s and 0s so that it can be stored and/or transmitted

Digitizing Analog Audio

- Create a number of samples (per unit time)
- For each sample, identify a discrete level, which is then represented digitally
- Levels
 - Represented as bits: usually 8-16 bits (256 – 65356 levels)
 - Sufficient to cover the range of human hearing
 - Fewer levels introduce more quantizing noise
- Sample frequency
 - Measure in “kHz”, thousands of samples per second

Uncompressed Digital Audio

- CD is a good example
 - Two channels (stereo)
 - 44,100 samples per second (44.1 kHz)
 - 16 bits per sample
- $2 * 44,100 * 16 = 1.4 \text{ Mbps}$
- Speech (or low quality “anything else”)
 - Mono, 8 kHz, 8 bits per sample
 - Equals 64 kbps
- All of this is **without compression**

Mu-Law Encoding

- Not much compression at all (so really just an encoding method), but technically there is some compression
- The compression that does exist is to covert 14-bits of sample value using only 8 bit samples

Mu-Law Encoding

- Not much compression at all (so really just an encoding method), but technically there is some compression
- The compression that does exist is to convert 14-bits of sample value using only 8 bit samples
 - Does this by having non-linear levels
 - Based on fact that humans are unable to differentiate very low and very high frequencies
 - So the quantizing noise generated by 8-bit non-linear encoding is the same as a 14-bit linear encoding
- Samples (8 bits) are independent and constant bit rate
 - Really useful for HW#2
- But if we don't want simple, how do we compress?

So Where Does Compression Happen?

- The key fact in audio is that from one sample instant to the next, the information is often the same (or very similar)
- So instead of just encoding the sampled audio value, the difference from the previous value is encoded
 - The difference between Pulse Code Modulation (PCM) and Adaptive Differential Pulse Code Modulation (ADPCM)
- Differential coding often results in long runs of very small to zero sample values
 - These runs can then be compressed using Run Length Coding
 - Other techniques like Huffman coding can be used
- These are all just simple data compression techniques

Next Generation Compression

- Take advantage of the way that the human ear perceives and prioritizes audio
 - We'll see similar ways in video compression that take advantage of the way the human eye perceives and prioritizes video
- Two **metrics** to consider
 - Quantitative: the difference between the *original* analog signal and the *reproduced* digital analog—measured as a noise value (SNR or dB)
 - Qualitative: several scales that are used as part of human testing

MP3

- MPEG (Motion Picture Experts Group), Part 3
 - Other two parts are for “system” (audio plus video) and video
 - Not to be confused with the fact that MP3 has three layers
- Wanted 6-to-1 compression
 - Original was ~1.5Mbps (2 channels, 44.1 kHz, 16 bit samples)
 - Goal was ~256 kbps
- “Auditory masking”
 - Strong signal drowns out weaker signal
 - Convert audio signal to frequency
 - Divide the frequencies into sub-bands
 - Digitize the sub-bands

Other Schemes

- Lots of work to develop different ways of collecting, encoding, and compressing samples
 - Transform the data
 - Filter out unimportant data
 - Re-sort data to achieve better compression
- Different schemes optimized for voice, music and other types of audio
 - Voice is surprisingly redundant
- Some protocols also incorporate aspects like security, resiliency, etc.

Tradeoffs to Consider

- The more complex the encoder, the more delay it adds
 - Once upon a time, encoder complexity could overwhelm the processing resources of the device, but not so much anymore
 - Only consideration might be battery power on mobile device
- The better the encoder and the greater the compression the more packet loss will affect the stream

Dealing with Loss

- Many compression schemes are not good at dealing with lost data (e.g., MP3)
 - If you cut data out of an MP3 file, it will likely crash the player
- Work to develop robust encoding schemes has become less important
 - Most streaming just uses **TCP** and progressive download
 - ...except video conferencing and VoIP
- **Techniques to deal with packet loss?**

Dealing with Loss

- Many compression schemes are not good at dealing with lost data (e.g., MP3)
 - If you cut data out of an MP3 file, it will likely crash the player
- Work to develop robust encoding schemes has become less important
 - Most streaming just uses TCP and progressive download
 - ...except video conferencing and VoIP
- Techniques to deal with packet loss
 - Use small packets (adds overhead)
 - Use loss concealment (e.g., play last packet)
 - Scramble and de-scramble the data (but introduces delay)
 - Add forward error correction (adds overhead)

HW#2 Overview

- Straightforward set of programming steps
 - Open source file
 - Read chunks (i.e., packets of data)
 - Decide if received or lost
 - Write data or if lost, write something in its place
 - Ending file size will be exactly the same
 - Open and listen to resulting file
- Use mu-law encoded packets (.au format)
 - Samples are 1 byte (read as unsigned ints) and independent
 - Can replace individual samples easily
 - Always copy the header
 - http://en.wikipedia.org/wiki/Au_file_format
- Run some experiments and report on the results
 - Packet size, loss rates, basic error concealment