

Jan-30 Lecture

Video Basics


- Represented by 3 8-bit values
 - RGB: Red, Green, Blue
 - YUV: Luma (Y) and two chrominance: Cr (red) and Cb (blue)
- HDTV as an example
 - $1920 \times 1080 \times 24 \text{ bpp} \times 30 \text{ fps} = 1.5 \text{ Gbps}$ (60 Hz interlaced)
- Other smaller formats
 - NTSC: 352×240 @ 30fps
 - PAL: Source Input Format (SIF): 352×288 @ 25 fps
 - Quarter Common Intermediate Format (QCIF): 176×144
 - Sub QCIF: 128×96
 - 4CIF: 704×576
 - 16CIF: 1408×1152

Video Basics

- Three mechanisms used as part of encoding/compression scheme
 1. **Spatial**: similarities around a given location of a frame
 2. **Temporal**: similarities around a given location across time
 3. **Lossy**: eliminate details not visible to the naked eye
 - Downsampling and manipulating the bit stream
- As we discuss different techniques, pay attention to those that increase **delay** and **loss effects**:
 - Requiring future data to encode current data
 - Computational complexity
 - Differential encoding

1. Spatial Redundancy

- Focused on similarities within a single image
 - Take advantage of the fact that most images in a video have similar values in nearby positions
- One option is to use **differential encoding**
 - Similar to delta compression
 - Assume that adjacent values in an image are the same and only encode the difference



2. Temporal Redundancy

- Check for similarities between video frames
 - Can look in the **reverse or forward** direction
 - Can also use **reverse and forward** direction
- If no other frames exist (single image) or no temporal redundancy is used, the frame is **intra-encoded** (sometimes called an I-frame)
 - Takes considerably more storage/bandwidth than frames that can take advantage of temporal redundancy

2. Temporal Redundancy

- Predicted images (P-frames) are based on other I-or P- frames
- Bi-directionally predicted images (B-frames) are based on combination of forward and backward prediction

2. Temporal Redundancy


- Predicted images (P-frames) are based on other I-or P- frames
 - Encoder does an **expanding ring search** to find image components (**motion compensation**)
 - How far from original location to look corresponds to how much **processing** is necessary and how much compression is had (key reason encoder is more complex)
- Bi-directionally predicted images (B-frames) are based on combination of forward and backward prediction

2. Temporal Redundancy


- Predicted images (P-frames) are based on other I-or P- frames
 - Encoder does an **expanding ring search** to find image components (**motion compensation**)
 - How far from original location to look corresponds to how much **processing** is necessary and how much compression is had (key reason encoder is more complex)
- Bi-directionally predicted images (B-frames) are based on combination of forward and backward prediction
 - If imagine component is in Location A now and Location C in the future. Half way between now and then, it should be in Location B (**interpolation**).
 - **Encode the “error”**: the difference between predicted Location B and where it actually is

2. Temporal Redundancy


- **Thresholds** are used for P and B frames
 - If there are enough differences (e.g., a scene change) such that an P or B frame would not result in any less data, then the frame is encoded as an I frame
- Typically use a pattern of I, P, and B frames
 - Ex: I B B P B B P B B ... I B B P B B P B B ... repeat
 - **Could encode all I frames** (motion JPEG)
- For **real-time video**, typically no B frames
 - B frames depend on future frames, can't encode and send until the future I frame is generated (**so adds delay**)
- For compressed and stored video, different I/P/B patterns can be tried
 - Nice tradeoff between processing and encoding efficiency



(a)



(b)




(c)

Dec 2442(Frame 2442)



- See Time 1:39 in video

Frame 2442(Dec 2442)




- Green = redundant with other frame
- Red = newly encoded

Frame 2369(Dec 2369)




- See Time 1:37 in video

Frame 2369(Dec 2369)



- Green = redundant with other frame
- Red = newly encoded

Frame 2369(Dec 2369)




- Arrows show interpolation

Frame 2380(Dec 2381)



- See Time 1:38 in video

Frame 2380(Dec 2381)



- Blue= redundant with other frame
- Red = newly encoded

3. Lossy Compression

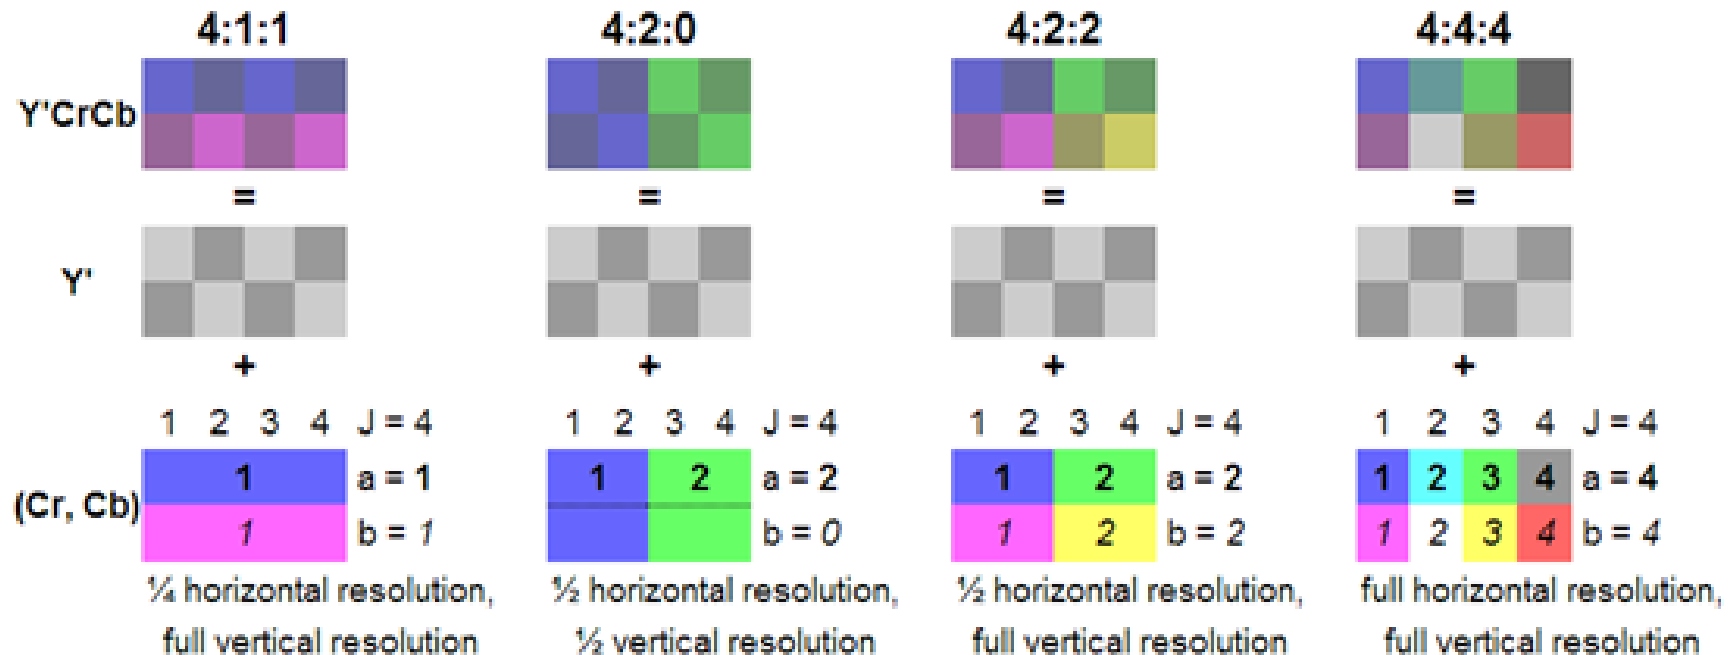
- a. Sub-sampling
 - Don't encode every bit of every block
 - This step adds loss
- b. Apply mathematical transforms
 - Usually something like a Discrete Cosine Transform
- c. Quantization
 - Eliminate differences in values
 - This step adds loss
- d. Stream coding
 - Use zig-zag and stream compression

3a. Sub-Sampling

- Eye is most sensitive to changes in luminance, and less sensitive to variations in chrominance
- Start with 16x16 block of pixels (macroblock)
 - Divide in 4 8x8 blocks
 - Full encoding is 12 blocks per macroblock (4 Luma, 4 Cr, and 4 Cb)
 - Syntax is: a:b:c (Luma:Cr:Cb)

3a. Sub-Sampling

- Eye is most sensitive to changes in **luminance**, and less sensitive to variations in **chrominance**
- Start with 16x16 block of pixels (macroblock)
 - Divide in 4 8x8 blocks
 - Full encoding is 12 blocks per macroblock (4 Luma, 4 Cr, and 4 Cb)
 - Syntax is: a:b:c (Luma:Cr:Cb)






- See Time 1:40 in video



- Y (luminance--brightness) values only



- U (chrominance #1) values only

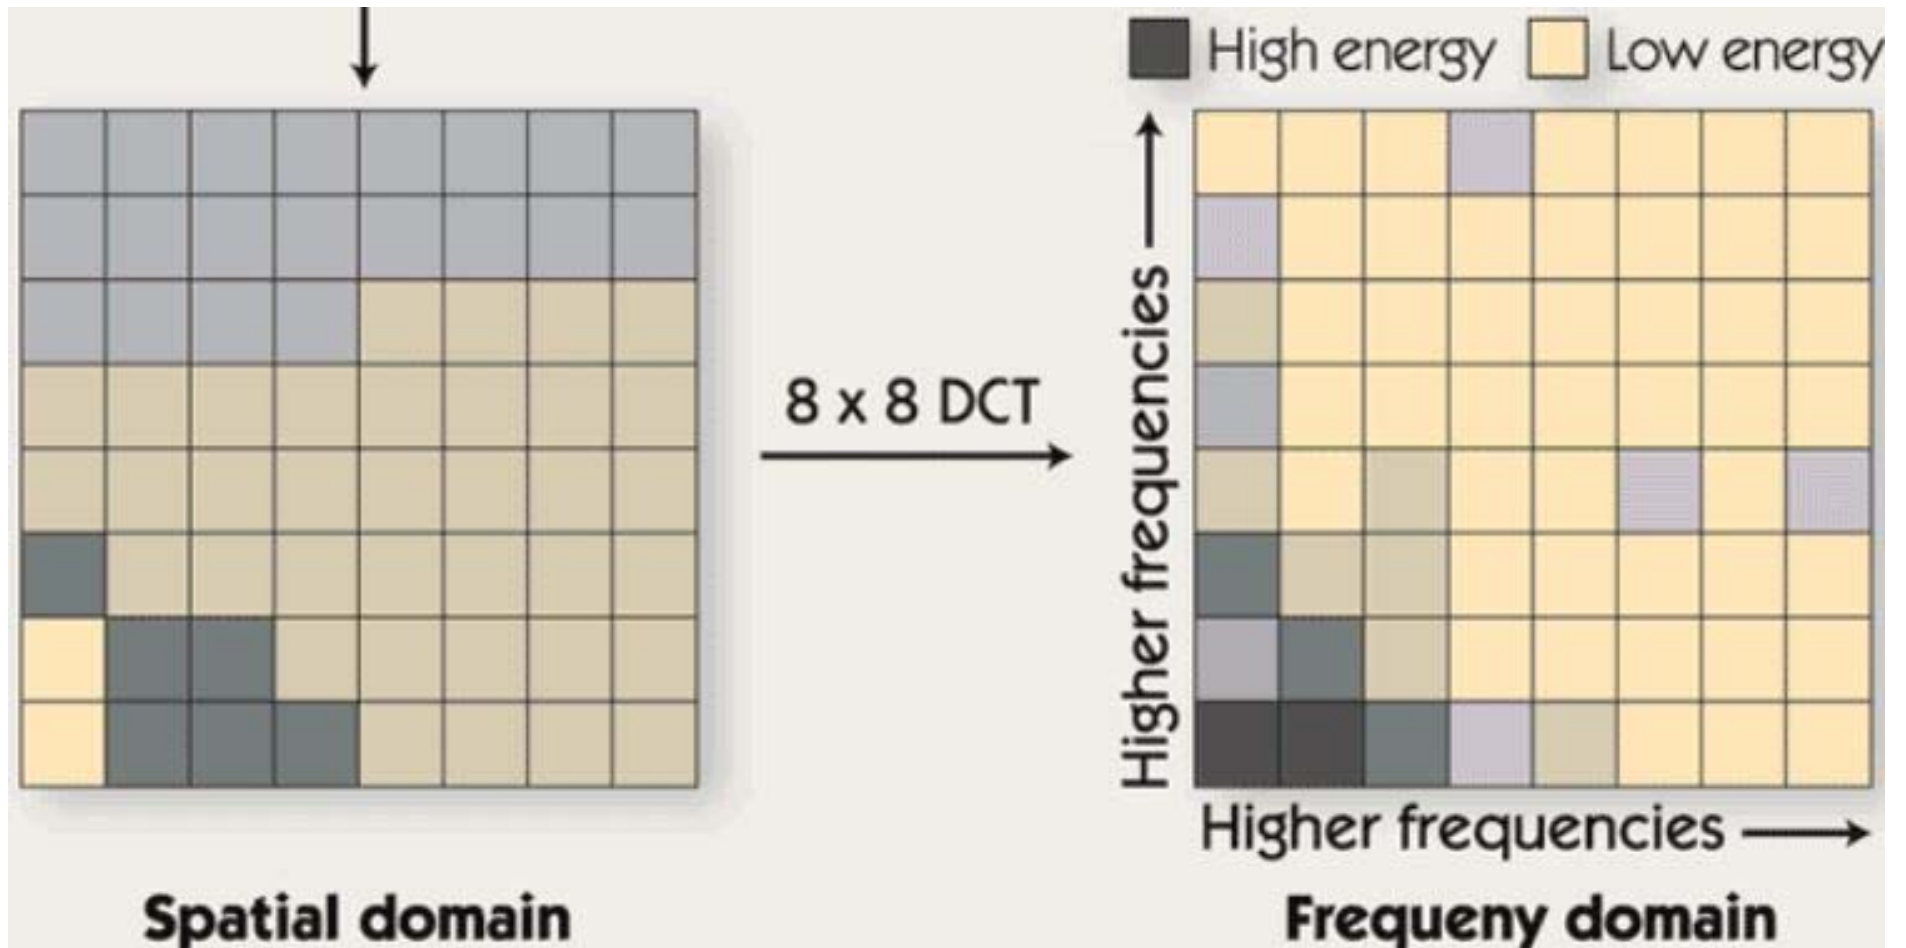


- V (chrominance #2) values only

3b. Transform Coding

- Transform coding is used to convert spatial image pixel values to **transform coefficient** values
 - No information is lost, the number of coefficients produced is equal to the number of pixels transformed
- The result is that most of the “energy” in the image will be contained in a few large transform coefficients
 - Generally, only a few coefficients will contain most of the energy in a block
 - Smaller coefficients can be **coarsely quantized** or deleted without doing visible damage to the reproduced image

3b. Transform Coding




3b. Transform Coding

- Many types of transforms have been tried for picture coding
 - Fourier, Karhonen-Loeve, Walsh-Hadamard, lapped orthogonal, discrete cosine, and wavelets
- Goals
 - The most concentration of energy
 - Least number of artifacts

3c. Quantization

- The level of quantization provides excellent tradeoff between quality and level of compression
 - More quantization means more compression which means **less bandwidth but more artifacts**
- Quantization can be adjusted dynamically
 - **Constant Bit Rate (CBR)**: same amount of bandwidth no matter the amount of energy/action in a picture
 - **Variable Bit Rate (VBR)**: bandwidth requirements vary based on complexity and motion in video
- Use of quantization is the source of noise/error in a compressed stream (different than network data loss)

3c. Error/Noise



3c. Error/Noise

- Coding error: the difference between the source picture and the reproduced picture
 - Just like for audio
- Coding error is measured as the root-mean-square between the two values
 - A common metric for evaluating the performance of an encoding system

3d. Stream Coding

- An example block of 8x8 DCT samples:

| | | | | | | | |
|----|----|---|----|---|---|---|---|
| 12 | 34 | 0 | 54 | 0 | 0 | 0 | 0 |
| 87 | 0 | 0 | 12 | 0 | 0 | 0 | 0 |
| 16 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

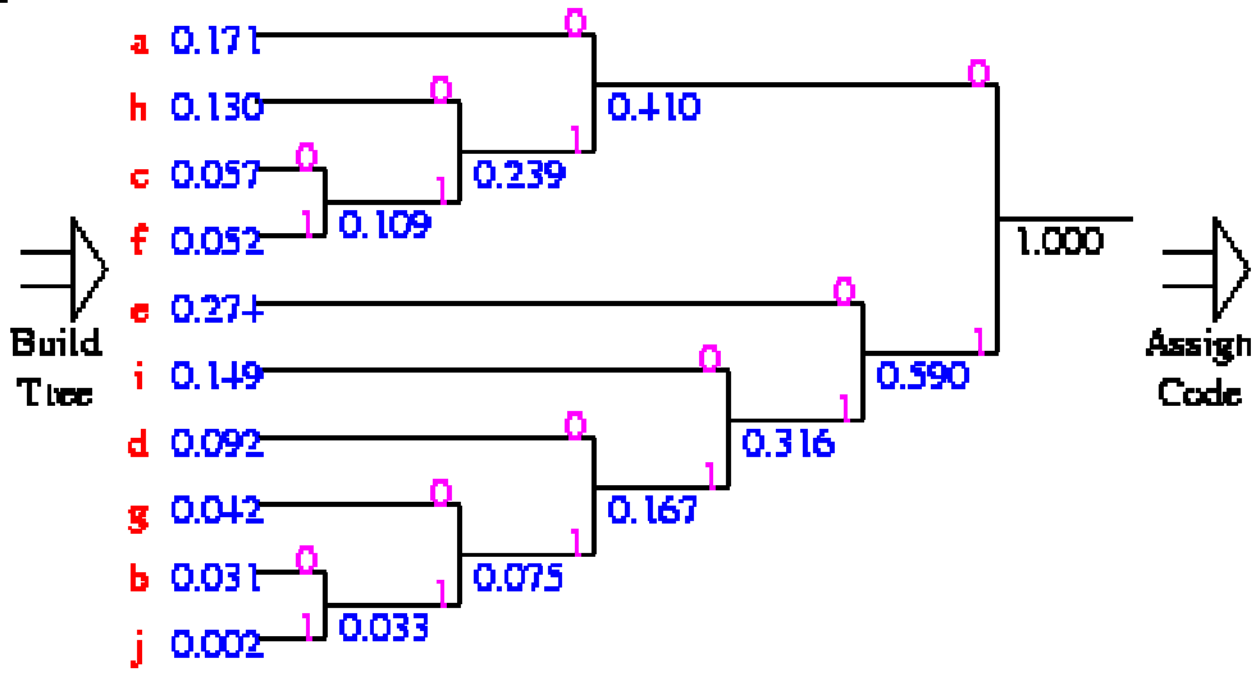
- First do zig-zag sequencing:
 - 12 34 87 16 0 0 54 0 0 0 0 0 0 12 0 0 0 0 0 0

3d. Stream Coding

- Zig-zag:
 - 12 34 87 16 0 0 54 0 0 0 0 0 0 12 0 0 0 0 0 0
- Then apply quantization:
 - 12 36 88 16 0 0 56 0 0 0 0 0 0 12 0 0 0 0 0 0
- Then apply run-length coding
 - Instead of long sequence of zeros, replace with “sequence of X zeros”
- Then apply Huffman coding
 - Most common sequence of bits is represented by shortest code

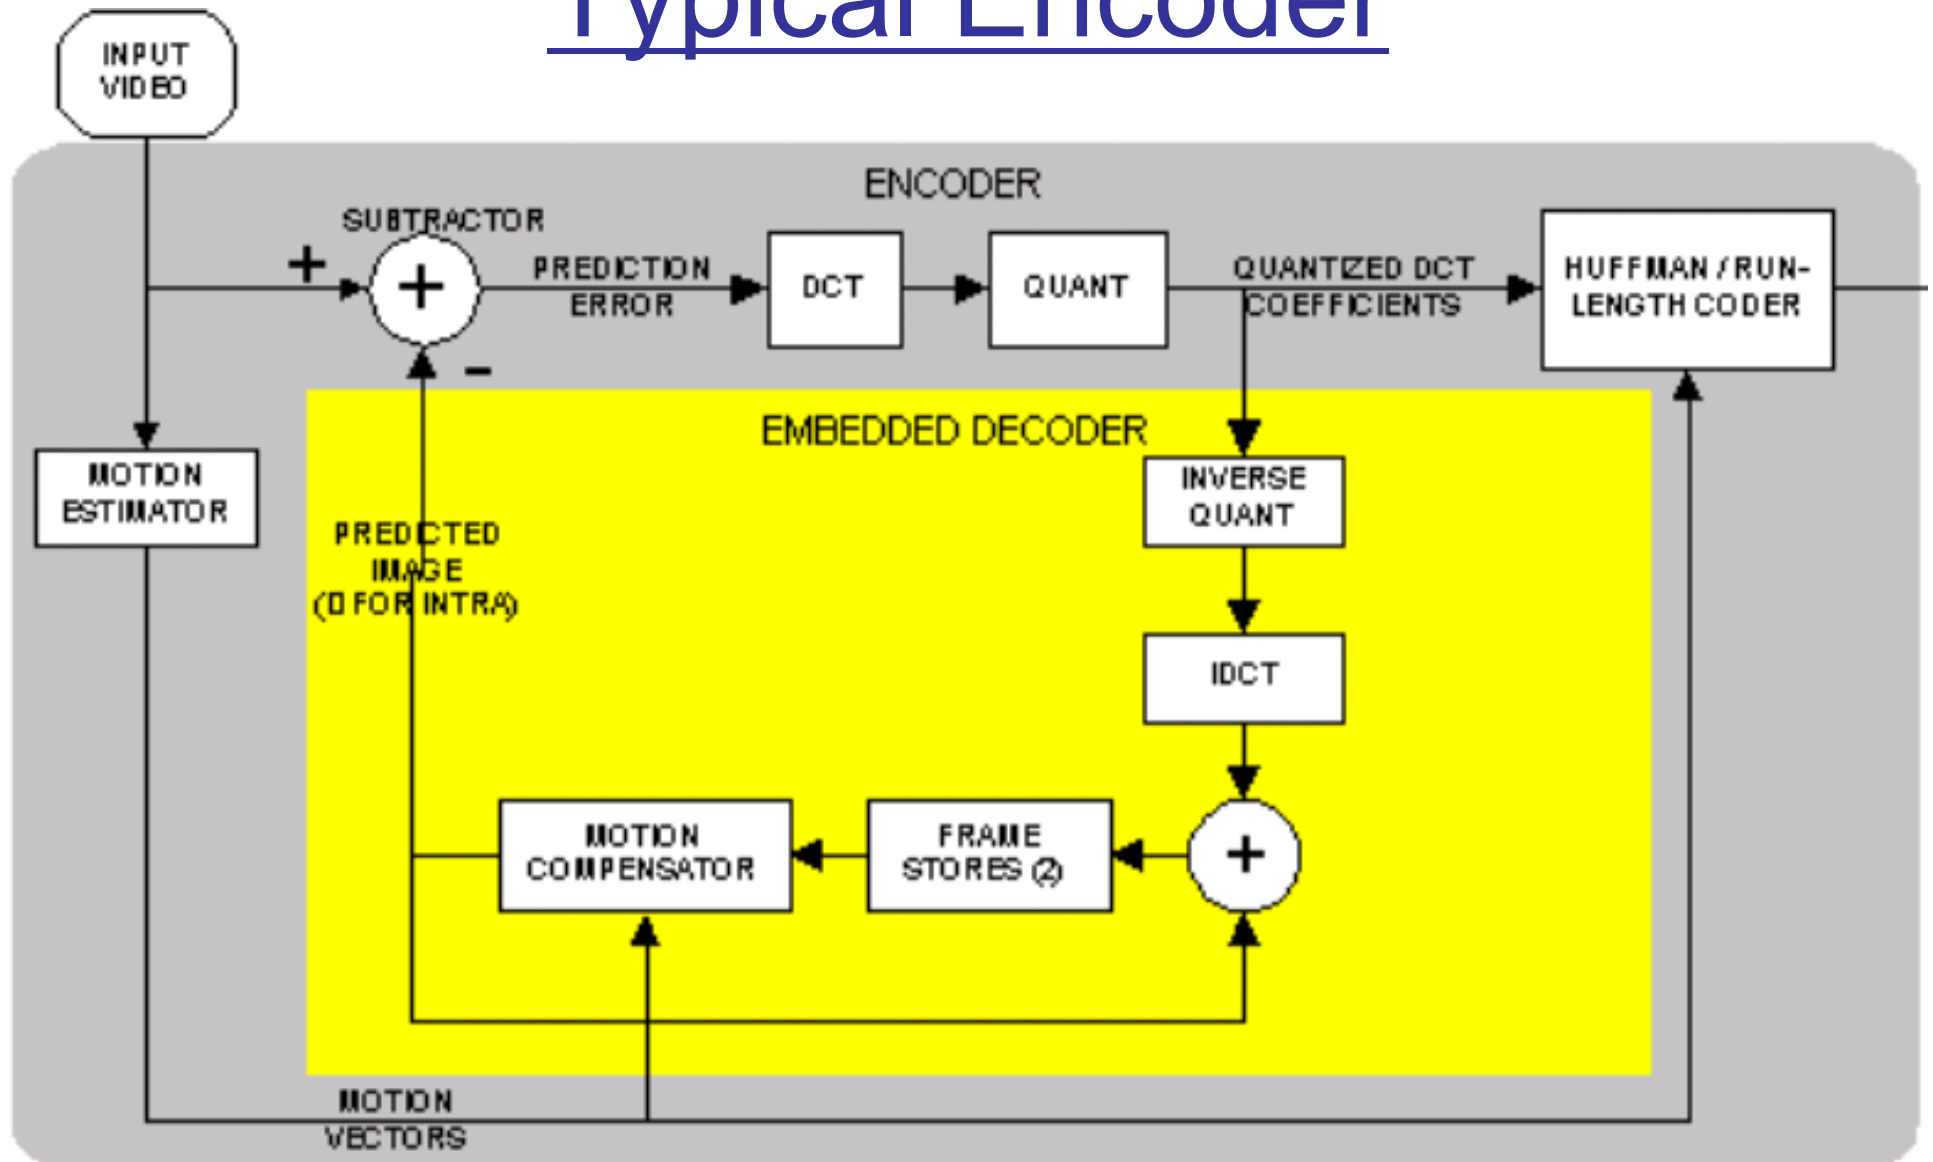
3d. Stream Coding

| X | p(X) |
|---|-------|
| a | 0.171 |
| b | 0.031 |
| c | 0.057 |
| d | 0.092 |
| e | 0.274 |
| f | 0.052 |
| g | 0.042 |
| h | 0.130 |
| i | 0.149 |
| j | 0.002 |



| X | Code |
|---|--------|
| a | 00 |
| b | 111110 |
| c | 0110 |
| d | 1110 |
| e | 10 |
| f | 0111 |
| g | 11110 |
| h | 010 |
| i | 110 |
| j | 111111 |

Typical Encoder



Standards

- MPEG-1 (1993)
 - Designed up to 1.5 Mbps
 - Standard CD-ROM, NTSC video quality
- MPEG-2 (1995)
 - Designed for between 1.5 and 15 Mbps
 - Standard for DVD, HDTV
- MPEG-4 (1999)
 - Object-based compression
- MPEG-7 (2002)
 - Provides framework for adding descriptive information about video contents: uses XML to store meta-data
- MPEG-21(2001)
 - Adds digital rights/permissions/restrictions

Other Standards

- MPEG came out of ISO
- Also CCITT (which became ITU)
 - Early on, principally designed encoding for low-bit rate video conferencing (Ex: H-261, H-263, H-264, etc.)
 - Typically use the same components (e.g., temporal, spatial, etc.) in the encoding scheme
- Typically there are separate standards for audio and video and then for a combination of the two

MPEG-1

- Consists of 5 parts:
 - Systems (storage and synchronization of video, audio, and other data together)
 - Video (compressed video content)
 - Audio (compressed audio content)
 - 3 different layers, the third is most commonly used (MP3)
 - Conformance testing (testing the correctness of implementations of the standard)
 - Reference software (example software showing how to encode and decode according to the standard)

MPEG-2

- Pictures are either “frame picture” or “field picture”
 - “frame picture”: complete frame
 - “field picture”: half of interlaced frame
- Has a header
 - Picture type (I, P, B)
 - Temporal reference information
 - Motion vector search range
 - Optional user data
- **Sequence** is a group of **Groups of Pictures (GOPs)**
 - Series of I, P, and B frames
- Further divided: **Slice** is a row of **Macroblocks** is group of **Blocks**

Other Standards

- Almost all other coding standards have the same basic elements
- Differences include:
 - Different techniques within each function
 - Ex: a different transform
 - Different “stream” structure (e.g., headers and data org)
 - lossy transmission
 - asymmetric coding: processing-intensive encode but easy decode
 - battery efficient decode
 - limited capacity transmission links
 - Some include special features
 - Ex: **Google's VP8** has a golden frame

Network Transmission

- Only some standards are designed for transmission
 - What happens when **part of an I frame or P frame** is lost?
 - What happens when **header information** is lost?
- There are steps that can be taken to improve resiliency
 - Basic idea is to **have fewer critical elements**
 - Not have multiple/many frames rely on same portions of the stream
 - Increased redundancy means less susceptible to loss, but at the tradeoff of sending more information

Network Considerations

- Most video now is TCP-based
 - Even for real-time and interactive streams
 - So use of UDP-based streaming is fairly limited
- Still, even with TCP, the goal is to avoid congestion and TCP backoffs
 - Catastrophic if TCP is used for interactive
 - Problematic if TCP is used for real-time
- A difference from audio is that video can consume a significant amount of bandwidth and congestion becomes a real concern

Congestion Solutions

- Rely on the user to choose the right quality
- Develop a dynamic solution that monitors loss rates and changes the quality in response to congestion (or the increasing likelihood of congestion)
 - Monitor throughput and increase/decrease quality
 - Increase/decrease frame rate
 - Increase/decrease quantization
 - Increase/decrease frame size
 - Use layered encoding

Scalable Video Coding

- Create set of *dependent* layers such that each additional layer adds detail and clarity
- Very useful for creating and sending a video stream to lots of people simultaneously but who have different end-to-end throughput levels
 - Adds complexity and inefficiency (more overhead to encode dependent layers)
- **Alternative** is to send separately encoded streams
 - Same stream encoded at different levels—requires switching to different stream

H.264 Scalable Video Coding

- One standardized solution uses H.264, an ITU standard for video compression
- Basic idea is how to do **scalable video coding**
 - Additional B frames can be added
 - Additional B frames are dependent on the higher level B frames directly on either side of the current level

Modern Video Streaming

- M. Tourad Diallo, H. Moustafa, H. Afifi, N. Marechal, “Adaptation of Audiovisual Contents and Their Delivery Means,” Communications of the ACM, vol. 56, no. 11, pp. 86-93, November 2013.