

Feb-20 Lecture

Putting Some of it Together

- Continue to probe the inter-workings of ISPs
 - Routing and peering
 - Route stability
 - IPSec then oh so briefly on firewalls
- We'll talk about additional services like content delivery next week

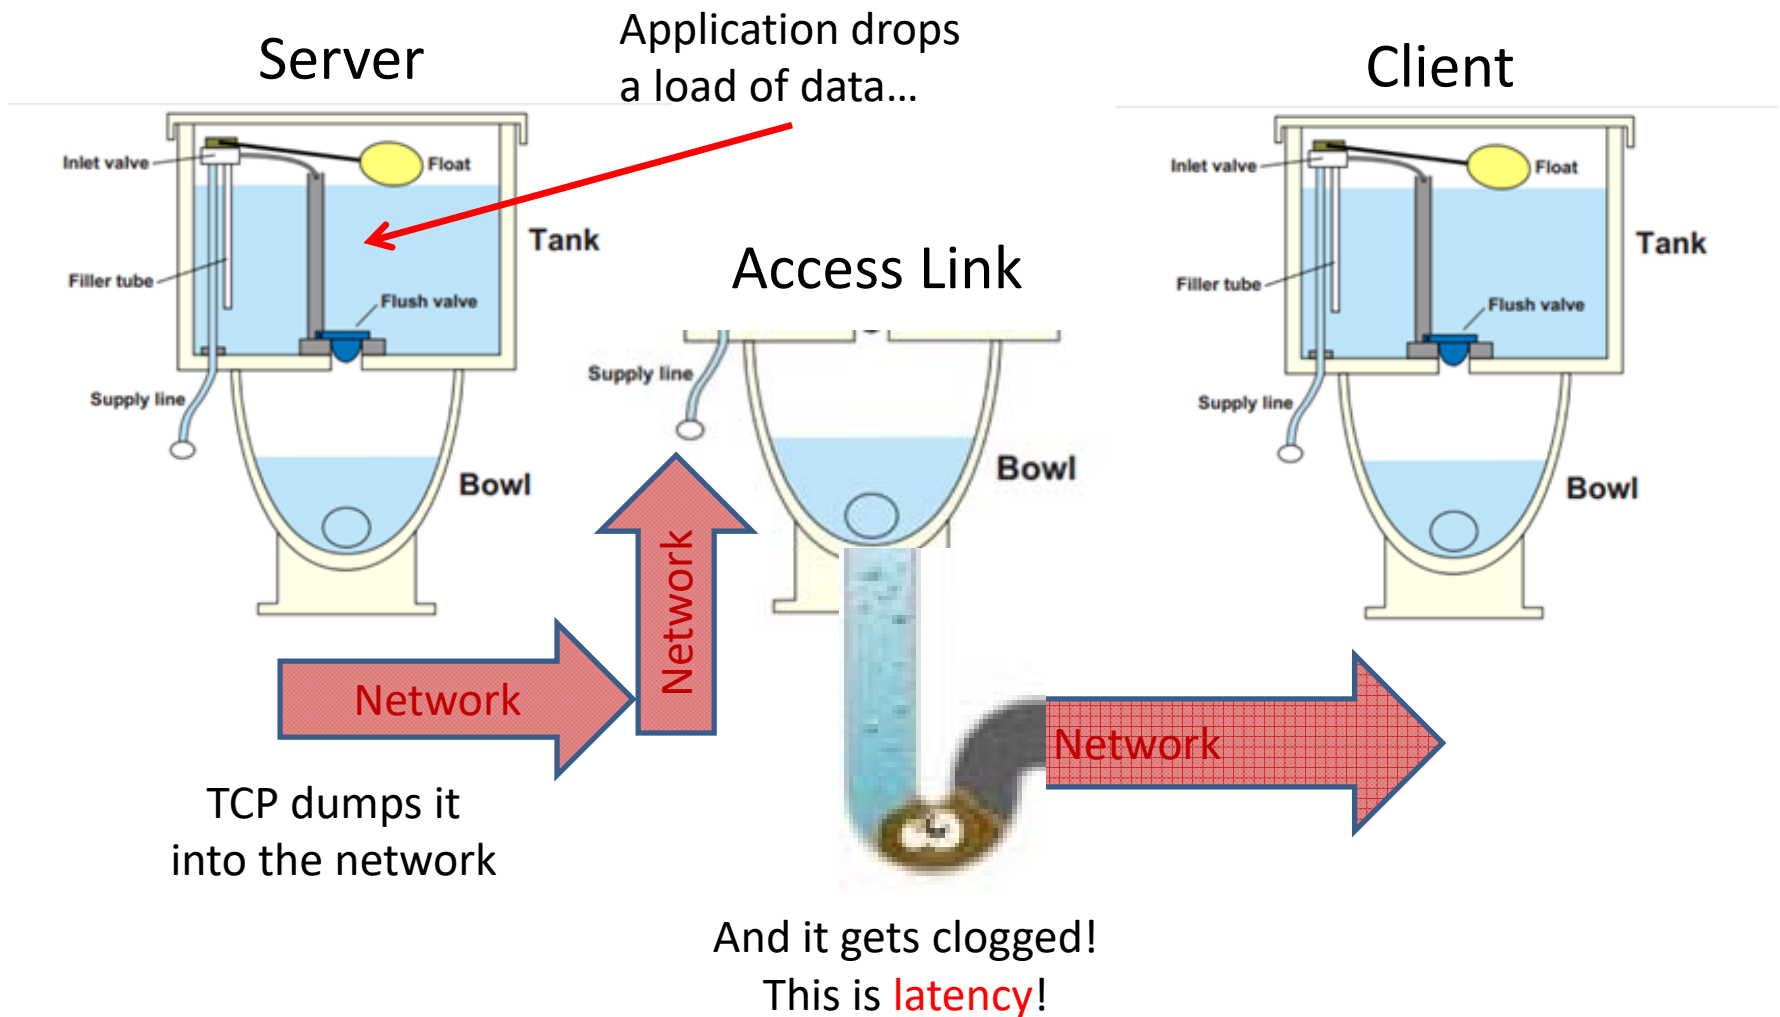
ISP/Core Impacts/Affects on TCP

- TCP behavior is heavily dependent on latency, jitter, and loss


ISP/Core Impacts/Affects on TCP

- TCP behavior is heavily dependent on latency, jitter, and loss
 - **Latency**: affects how quickly TCP gets through a “round” and therefore, add to its cong window
 - **Jitter**: also affects how responsive/accurate TCP can be in its timeout
 - **Loss**: more dramatic affect on cong window, but depends on the type of cong control strategy

The Source of Latency

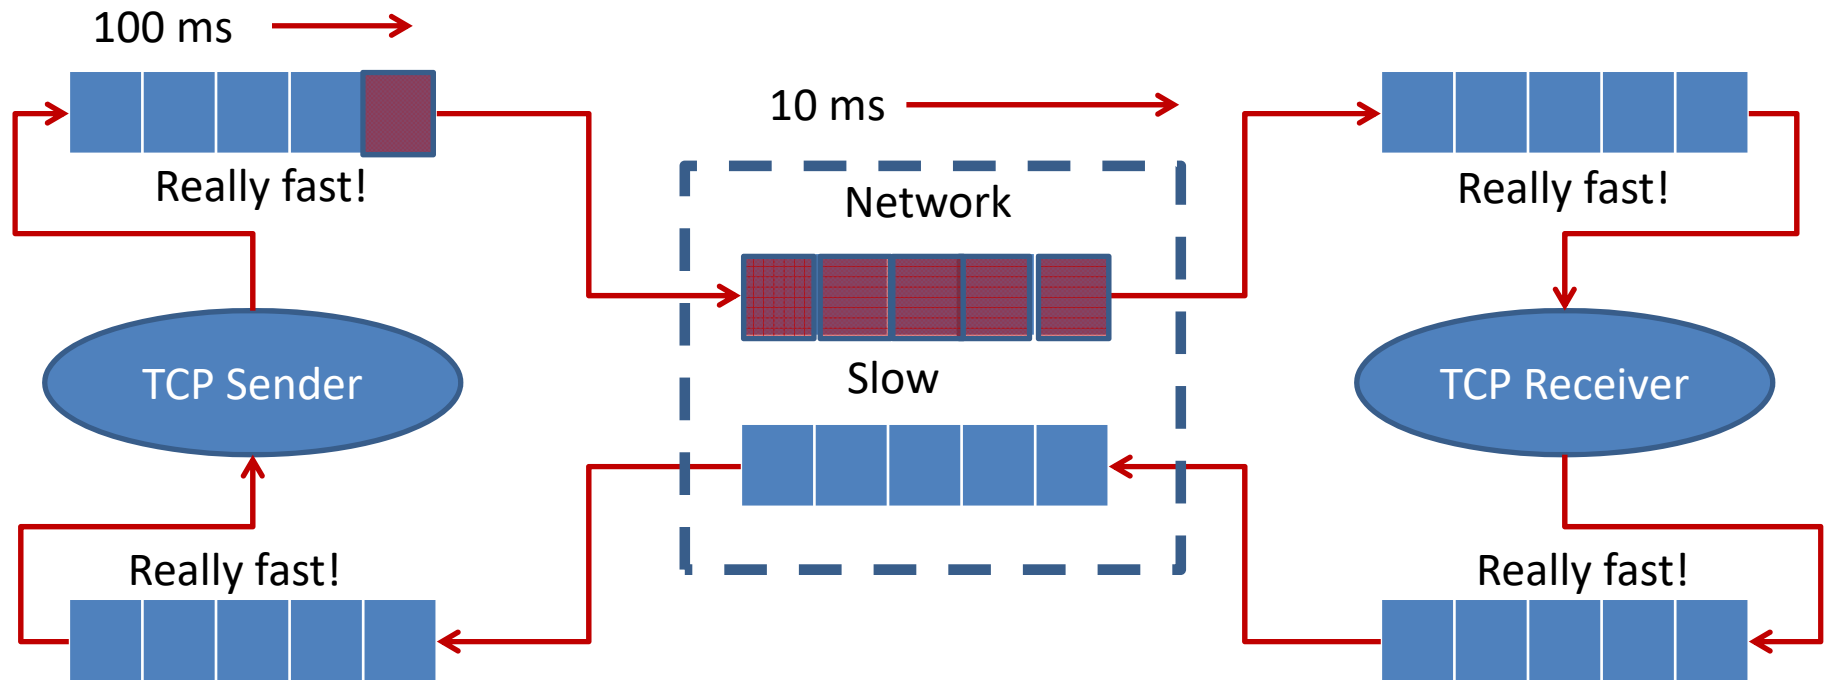


TCP Control Loop



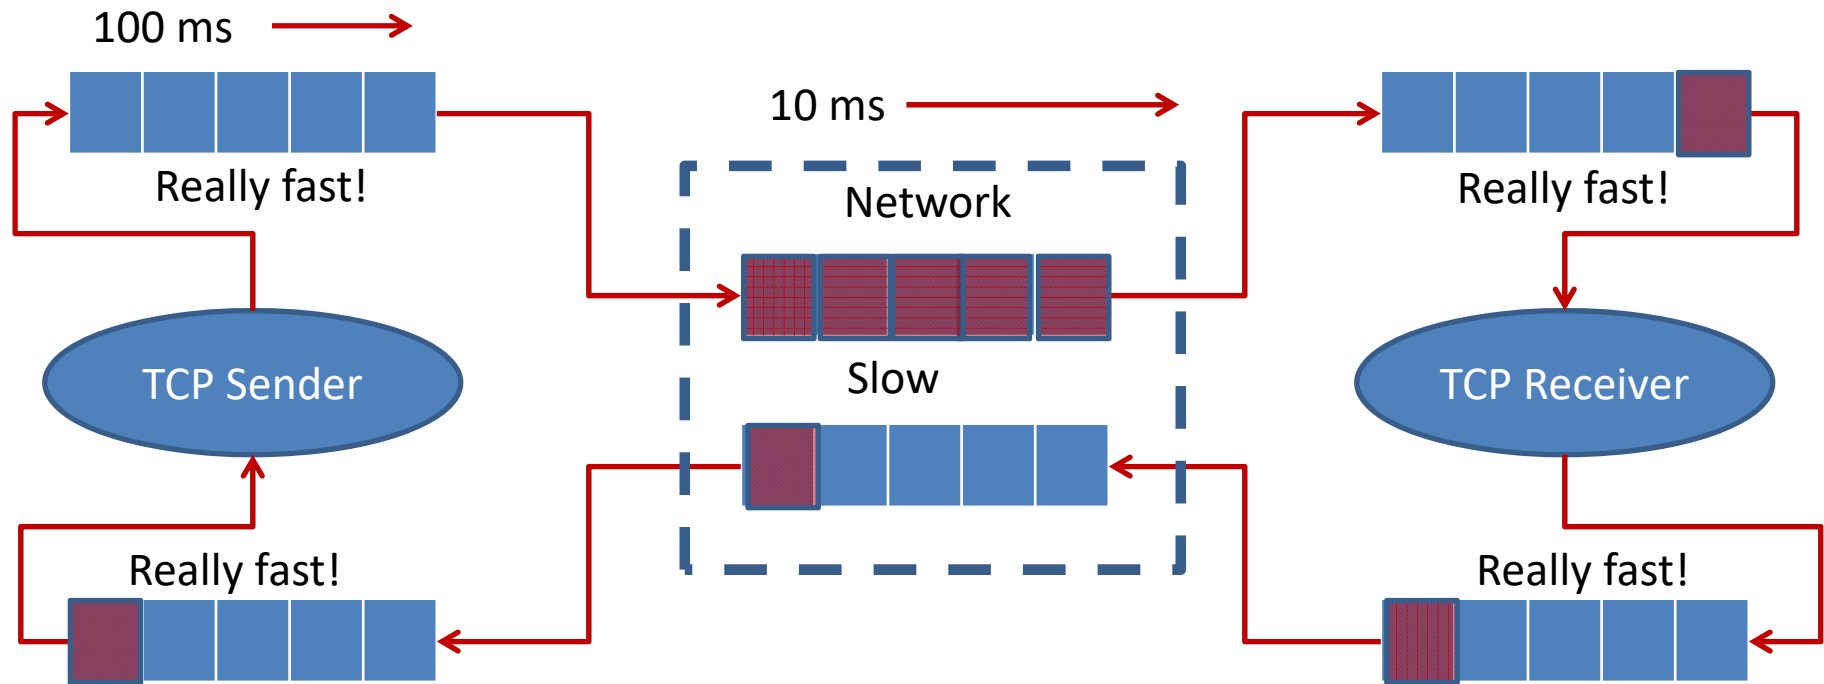
- Using **Congestion Control** and **ACK clocking** TCP will:
 - Calculate a SND.WND size that sends enough packets to keep the network busy for an RTT

TCP Control Loop




- Using **Congestion Control** and **ACK clocking** TCP will:
 - Calculate a `SND.WND` size that sends enough packets to keep the network busy for an RTT

TCP Control Loop



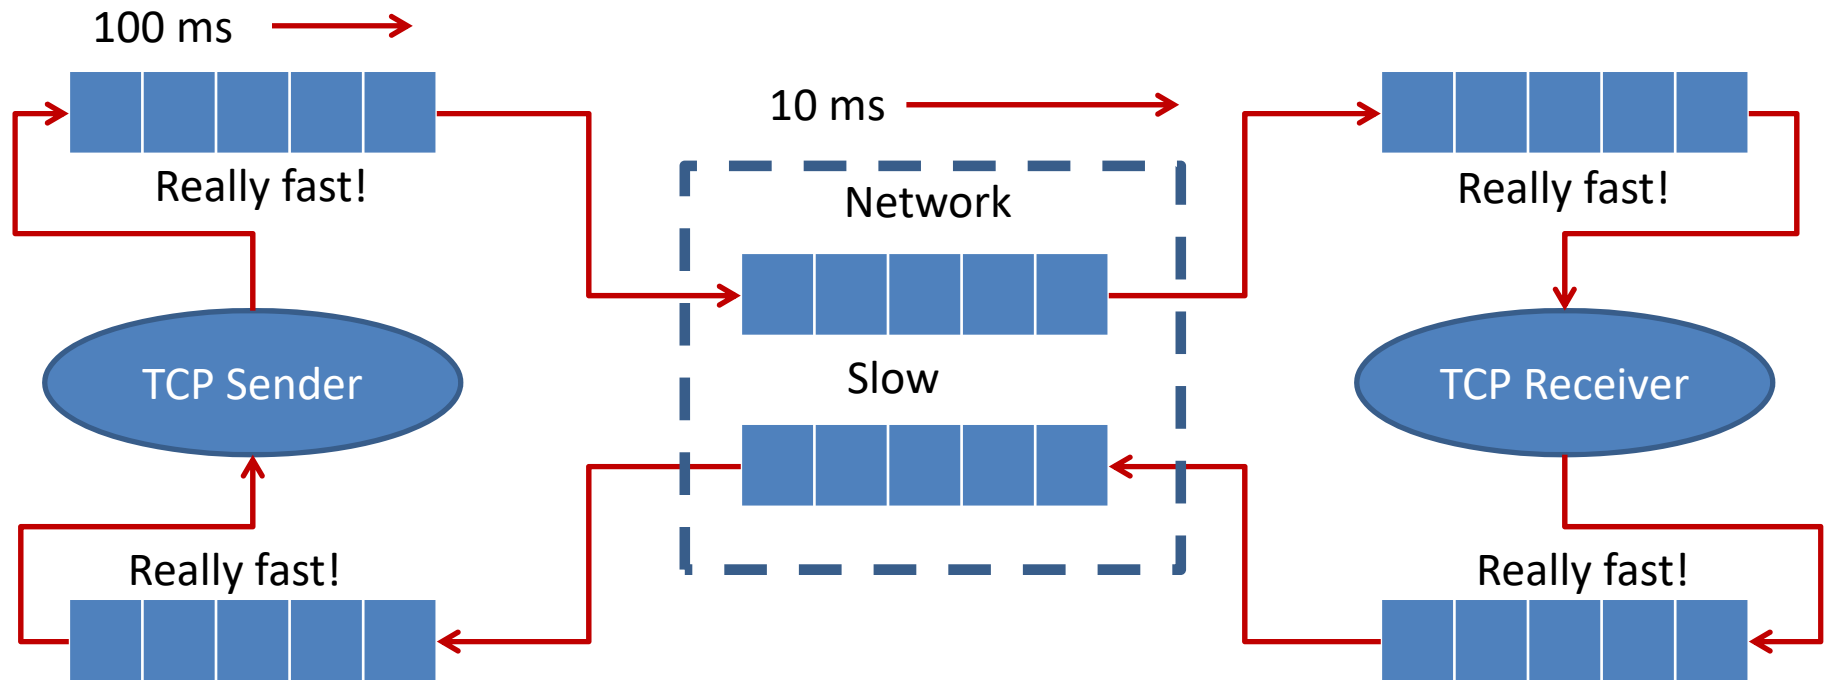
- Using **Congestion Control** and **ACK clocking** TCP will:
 - Calculate a `SND.WND` size that sends enough packets to keep the network busy for an RTT

TCP Control Loop



- By the time the first ACK arrives at the sender...
 - The sender queue will have drained
 - That's the goal at least


TCP Control Loop



- When packet loss does occur, it happens in one of two places:
 - On the ingress (packets arrive faster than can be processed)
 - On the egress (too many packets trying to go the same place(s))


What Happens Next?

- Packets pile up at the slowest device (queue)
 - All the other queues are faster and will drain
 - Queue management must be applied at the slowest queue
- Why?



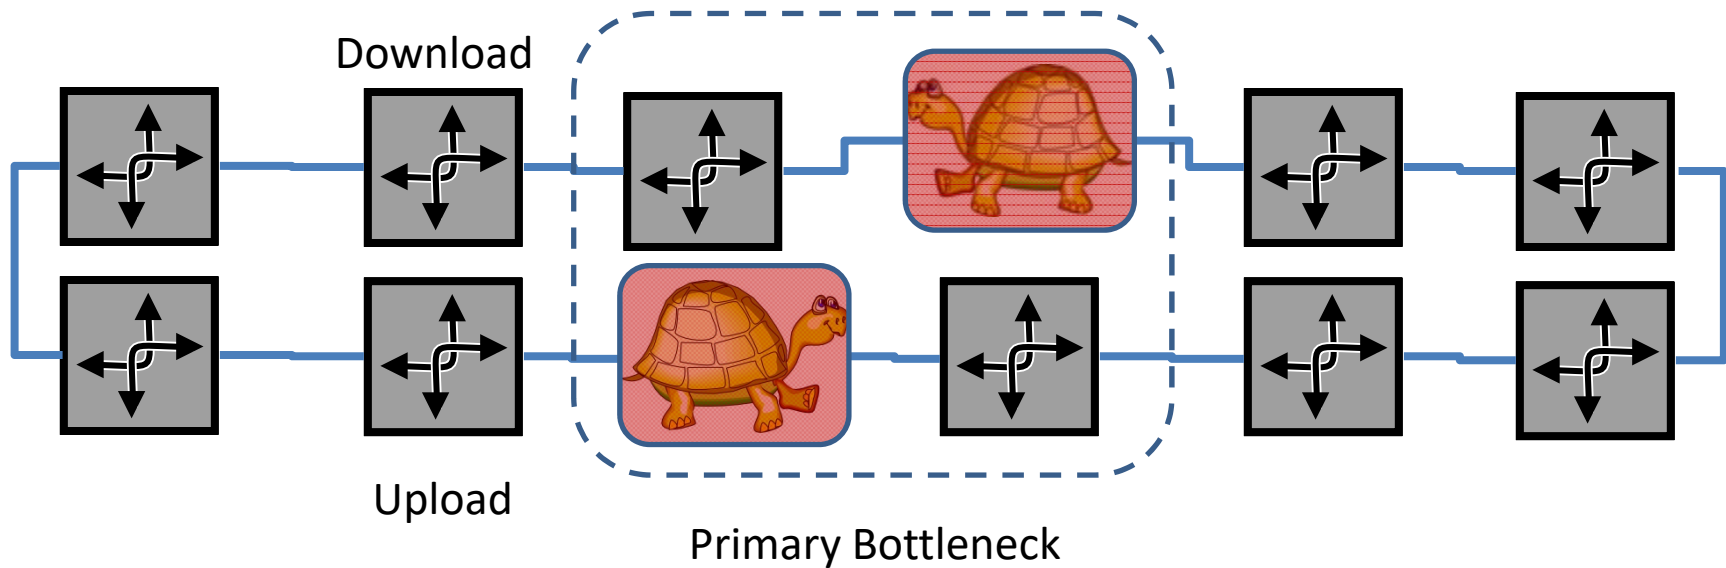
What is Bufferbloat?

- Memory is cheap
 - If packets are being lost on ingress/egress, just add more memory/buffer
 - ...and more and more and more
- What are the consequences?




How To Solve Bufferbloat?

- Instead of trying to prevent packet loss, allow it to happen
 - TCP is good at adjusting
- Better yet, use packet loss to “signal” to TCP that congestion is occurring




Random Early Detection (RED)




- When queue grows beyond “low water mark” start to drop packets with $\text{Prob} = p$ (low p)
 - Can drop or just mark them

Random Early Detection (RED)




- As queue fills, become more aggressive on dropping probability
- When queue is full, nothing left to do, but drop everything

Random Early Detection (RED)



- As always, there's a challenge!
 - What to set queue thresholds to?
 - What to set p values at?

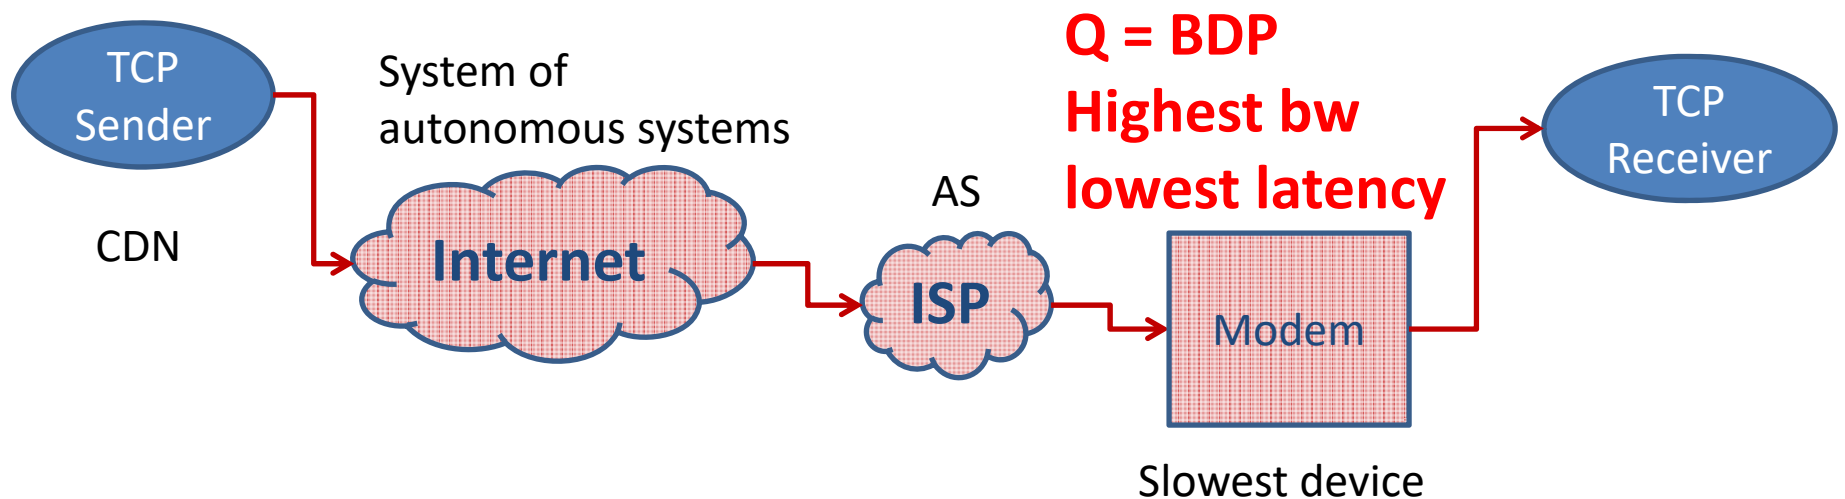
Random Early Detection (RED)



- The answer is based on the **Bandwidth Delay Product**

Bandwidth Delay Product (BDP)

- TCP assumptions:
 - Queue size = BDP
 - Highest bandwidth
 - Lowest latency
- What is BDP?
 - Bandwidth = The **output rate** of the **slowest device** in the path
 - Delay = Round Trip Time (RTT)



Queue Size Equations

- Queue Size = Bandwidth * Delay
– Bytes = bytes / sec * sec

Queue Size Equations

- Queue Size = Bandwidth * Delay
 - Bytes = bytes / sec * sec
- $Q = \text{BDP}$
 - If $Q = \text{BDP} \rightarrow$ full bandwidth, lowest latency
 - If $Q < \text{BDP} \rightarrow$ lose bandwidth, lowest latency
 - If $Q > \text{BDP} \rightarrow$ full bandwidth, high latency (Bufferbloat)
- However, we do not know the **bandwidth** or the **delay** in advance.
 - So how do we set the queue size?

How do you configure RED?

- **Yes**, but tuning is supposed to be difficult...

1. Red parameters:

- 2. **Bandwidth** 2 Mbps
- 3. **Delay** 200 ms
- 4. Avpkt 1000 bytes
- 5. Limit 8 * max

- 1. Min = BDP
- 2. Max = 2 * Min
- 3. Burst = (min+min+max) / (3 * avpkt)
- 4. Limit = 8 * Max
- 5. Min 50,000 Bytes
- 6. Max 100,000 Bytes
- 7. Burst 67 segments
- 8. Limit 800,000 Bytes

- Just set some “reasonable” values and see how it works

Some Experiments

- When you plug in numbers, the queue size is a simple calculation (and fixed)
- $Q = BD$
- $Q = 2 \text{ Mbps} * 200 \text{ ms}$
 - $Q = 400,000 \text{ bits (or 50kB)}$
- But what happens if the selected queue size isn't right?

Some Experiments

- When you plug in numbers, the queue size is a simple calculation (and fixed)
- $Q = BD$
- $Q = 2 \text{ Mbps} * 200 \text{ ms}$
 - $Q = 400,000 \text{ bits (or 50kB)}$
- But what happens if the selected queue size isn't right?

Queue Size Equations

- $Q = \text{BDP}$
 - If $Q = \text{BDP}$ \rightarrow full bandwidth, lowest latency
 - If $Q < \text{BDP}$ \rightarrow lose bandwidth, lowest latency
 - Queue is too small and more packets are dropped than it should be
 - TCP responds by reducing how much data is sent
 - If $Q > \text{BDP}$ \rightarrow full bandwidth, high latency **Bufferbloat**
 - Queue is too large
 - TCP doesn't see any drops
 - But RTT grows to unreasonable lengths

Next Steps

- RED was one of the first solutions
- Many others since RED
- They create complex mathematical models that can use network information to estimate queue size
 - Packet pairs, ping times, statistical distributions

The Evolution Continues

- Where once ISPs were just conduits for traffic, now there is much more engineering
- Early efforts were abysmal failures
 - See, e.g., integrated services (int-serv), the Resource Reservation Protocol (RSVP), and ATM
- More recent efforts have taken lessons from the past and been less processing-intensive
 - MultiProtocol Label Switching (MPLS)
 - Now Software Defined Networking